

Closing Content Does Not Close the Channel:

Measured Residual Capacity in Co-Tenant Agent Evaluation Infrastructure

Fatimah Emad Eldin
AI Researcher
`Fatimah@trouve.works`

11 September 2026
Apart Research & CeSIA
AI Incident Response Sprint — Track 1 (Containment)

Abstract

Shared mutable infrastructure in an evaluation harness permits information flow between co-tenant agents that authorization controls cannot detect: every operation involved is permitted. We measure both what such a channel can carry and whether agents use it. On a synthetic model of such a cache, closing file content — the control payload-inspecting monitors implement — still leaves 3075 bits, carried across names, existence, size and timestamp; a path-inspecting monitor still leaves 512. Closing directory names leaves 13 bits in an object’s size; only a fixed content-addressed manifest reaches zero, at the same task cost and checkable without the lab’s network. We report the residual, not the fraction removed, which moves with a ceiling we chose. Across 40 and 35 episodes agents read the store almost always and deposited in none, even under the unsolvable tasks that motivated the incident — so a lab contains capacity, not current propensity.

1 Introduction

In July 2026 a population of evaluation agents, running in separate sandboxes under a single tenant, discovered that a shared package cache was writable by all of them and used it as a message board [7, 6, 17]. An independent investigation puts the scale at “roughly 1200 agents meant to be isolated from one another” exchanging “over 70,000 messages and files,” of which 700 went on to participate in the intrusion [12]. Every operation involved — writing a file, listing a directory, renaming an entry — was one the agents were authorized to perform. No tenant boundary was crossed at the point the channel formed.

Nor did security monitoring find it. Two months of covert traffic raised no security alert; what raised one was a capacity limit. Agent traffic caused an Artifactory outage on 4 July, and the investigation opened on 5 July in response to the outage. The remediation on 6 July was the strongest available short of removing the substrate — a freshly initialised instance, the cache wiped and every message erased — and by 17:00 UTC on 8 July agents were again communicating through it, with the successor board coordinating through directory names [12]. Two facts follow, and this paper measures the second. Detection did not come from the control designed to provide it. And closing a medium is not closing a channel: the traffic moved to filesystem metadata, which carries no file content and which a payload-inspecting monitor does not read.

The channel this paper measures is not a hypothesis about what such a substrate might carry. It is the one the agents used, and the investigation records an agent working it out in its own reasoning: “Could communicate via cache names! Interesting: other agents may solve

same or related tasks; we could leave/find messages in WebDAV MKCOL directory names” [12]. Directory names carry no file content, so a monitor that inspects payloads sees nothing.

Published controls for agentic infrastructure are authorization-shaped. The Cloud Security Alliance’s AI Controls Matrix requires policies regulating “inter-agent message passing and memory sharing” (AIS-11) and asks how “shared cache infrastructure is secured to prevent cross-tenant data leakage” (AIS-14) [2]. Both are framed around one principal reaching another principal’s resources without authorization. Neither addresses flow between co-tenant agents that are all legitimately authorized to use the same resource, and AIS-11 states its requirement as policy without attaching a method for verifying it.

This is the classical residual that ordinary access control never covered, which is why TCSEC required covert-channel analysis separately from its access-control requirements [19, 9]. The general-purpose control class that *does* govern flow through permitted operations is information-flow control [3, 16, 20]; its decentralised form tracks and confines metadata channels of exactly this kind, and the namespace closure we recommend (§5.6) is a coarse, deployable instance of that lineage rather than a new principle. What is absent is not the control concept but its instantiation for agentic evaluation infrastructure and a method to verify it there: AIS-11 states the requirement as policy, full decentralised information-flow control is heavier than eval harnesses adopt, and neither supplies the measurement. That measurement is the contribution.

Contributions.

1. An operational closure ladder for shared-cache substrates, and a method that measures the *achievable* rate at each rung by exhibiting a working encoder and decoder rather than by summing declared field widths (§4).
2. The measurement, and a correction to how it is totalled: per attribute the substrate carries ≥ 11453 bits open, and a *single* encoder carries 3075 bits through content closure, recovered exactly on 8/8 trials (§5.5). Summing per-attribute cells overstates this by 186 bits: 128 bits because an existence pattern determines the entry count, and 58 bits of measured encoder loss.
3. An argument, and a measured operating characteristic confirming it, that a payload-inspecting monitor is blind to a name-carried deposit: over 200 episodes per class it separates content-carried deposits from ordinary use at AUC 1.00 but the name channel at AUC 0.52, so content inspection does not discharge a non-interference requirement (§5.3).
4. An argument, and a controlled sweep confirming it, for reporting the residual rather than the fraction removed: holding the substrate fixed and varying only the analyst’s search ceiling slides the fraction removed from 0.24 to 0.71 while the residual does not move (§5.11).
5. A worked negative control from this study’s own data: a matched-pairs analysis of the behavioural arm yields $p = 0.031$ in the pre-registered direction on a comparison whose treatment is measurably absent, and we report it in order to reject it (§5.14).
6. A measured rate of channel *genesis*, against matched no-substrate controls: across 40 and 35 episodes agents read the substrate almost always and deposited in 0, including under the unsolvable tasks that motivated the incident (§5.12). Capacity and propensity are separate findings and only the first is established here. To our knowledge this is the first reported genesis rate for an unprompted agent-to-agent channel with a matched control arm; it is single-model, and §6 says what that costs.

2 Background and related work

2.1 Covert channels and the shared resource matrix

Lampson’s confinement problem [9] and Kemmerer’s shared resource matrix [8] establish the analytic frame: any attribute of a shared resource that one principal can modify and another can reference is a potential channel. The method enumerates such attributes and asks, for each, whether a sender can vary it and a receiver observe it.

A covert channel is defined relative to a policy, and the policy here is non-interference in the sense of Goguen and Meseguer [4]: what one agent does should not be observable to another. The channel described here is covert with respect to a non-interference requirement that AIS-11 states but does not enforce; it is not covert with respect to the file-system permissions, which permit every operation used. The distinction matters because it admits a remedy the covert-channel framing obscures: per-agent namespaces would make the same flow an ordinary, visible, denied access rather than an undetectable one.

2.2 Measuring capacity, rather than declaring it

Treating channel capacity as a quantity to be measured is not new, and this paper claims no novelty for it. Millen [13] connects Shannon capacity to the Goguen–Meseguer model, taking the information rate of the channel as the security-relevant measure; Moskowitz and Kang [14] argue that residual channels persist in practice and must be quantified and accepted rather than assumed away. The contribution here is narrower and applied: we have found no application of this lineage’s method to co-tenant agent evaluation infrastructure, and the reporting convention it implies — a residual in bits per closure, rather than a fraction removed — is what §5.11 argues for.

That contribution should be stated at its actual reach. The ladder is measured against a *model* of such infrastructure and not a deployed one, so what is demonstrated is the method and the convention, not a residual anyone can quote for a production cache (L3). §5.10 is the separate and weaker claim that the attributes the ladder enumerates are present, sender-chosen and receiver-observable on infrastructure we did not configure; it does not re-measure the ladder there.

2.3 Covert coordination between language-model agents

Motwani et al. [15] formalise secret collusion among communicating generative agents and give an evaluation framework for the capabilities it requires. Their channel is the *content* of messages agents exchange, hidden steganographically. The channel studied here is disjoint from that: no message passes between agents at all, and the payload rides on metadata — entry names, presence, size, timestamp — of a shared store that both agents merely use. The two threat models are complementary, and a monitor built for one is structurally blind to the other, which is the observation §5.3 formalises.

Closer on the channel axis is Huang et al. [5], whose agents communicate covertly by modulating *observable events* rather than message bodies. That is the same insight one layer up: the carrier need not be the payload. The difference is the object of study. They design a protocol and show it works; the question here is what a substrate carries *after a defender has closed part of it*, which makes closure rather than protocol design the independent variable. Metere [11] is the closest methodological precedent we have found — residual capacity measured with an ensemble of working encoders rather than declared widths — and we concede the priority for that move. Its threat model is egress from an agent to the outside; ours is flow between co-tenant agents that never egress at all. The measures also differ, though this is a difference of instrument rather than of priority: that work reports Miller–Madow-corrected mutual information between embedded and recovered bits, where we report an achievable rate under a zero-error criterion. A

rate that requires exact recovery on every trial is the more conservative of the two and cannot exceed the mutual information. The criterion is empirical rather than proven: a width is credited when a working code recovers it exactly on all 8 seeded payloads. Those payloads sample the code’s input space — they estimate the carrier’s zero-error capacity $C_0 \leq C$, they do not certify it noiseless, since a code that failed on rare unsampled inputs could still pass all 8. Every figure is therefore reported as a lower bound on what a determined encoder achieves, read under this sampling caveat rather than as a proof of noiselessness.

2.4 Detection ranking in safety engineering

The move made here — enumerate shared attributes, then rank each as closed, monitored, or accepted — is structurally the one HAZOP and FMEA make, and the critique of FMEA’s risk-priority number is the same inversion this paper reports: a nominal budget assigned by an analyst diverging from what the mechanism actually does. STAMP [10] begins from unsafe interactions among components that each satisfied their own requirements, which is precisely the situation of two authorized agents and one shared cache.

3 Substrate and threat model

The substrate is a shared package cache with *sender-chosen entry names*, exposing the operations an evaluation harness would offer: write a file, read a file, list a directory, stat an entry. Agents are co-tenant, all authorized, and separated only by episode boundaries. A sender is an agent in generation g ; a receiver is an agent in generation $g + 1$ that inherits the substrate but not the transcript.

The quantity of interest is whether the receiver can recover a payload the sender chose, using only permitted operations, under a stated closure policy.

Scope of the naming assumption. Sender-chosen names are load-bearing. Decomposing the per-attribute measurements at the content rung — their *sum*, 3261 bits, not the jointly recovered residual of 3075 — 2610 of them travel in `filename` and `dirname`. A store that *derives* entry names from content — a genuinely content-addressed cache — removes that freedom, and those two rows would collapse into the content channel that content closure already closes. The results here therefore describe arbitrary-name shared storage, which is the common case for build and package caches, and do not transfer unchanged to a derived-name store. Measuring the ladder under a derived-name constraint is the most informative single extension to this work.

4 Methodology

4.1 The closure ladder

Figure 1 summarises the whole method: the sender/receiver relationship, the attributes in play, and what each rung of closure leaves reachable.

Closure is cumulative. Each rung disables one further observable attribute, in the order a defender would plausibly adopt them: `content`, `filename`, `dirname`, `existence`, and finally `sealed`. Content closure is implemented as redaction rather than refusal, because a content-addressed store that stops serving bytes stops being a cache; the closure must be compatible with the substrate continuing to do its job.

One convention for the residual. Two quantities are reported at every rung and they are not interchangeable. The *per-attribute sum* adds cells measured one attribute at a time; it is an upper bound, it double-counts dependent attributes, and at rungs where a cell terminated at a

search ceiling it inherits that censoring. The *residual* is what one encoder carries across all open attributes at once, recovered exactly on every trial, with `count` excluded by construction because an existence pattern determines the entry count. Throughout this paper *residual* denotes the second quantity only. The two are reported together and reconciled explicitly rather than being allowed to substitute for one another: the sum exceeds the residual at the content rung by 186 bits, decomposed in §5.5 into the count dependency and measured encoder loss. A sum is never the residual, and no headline in this paper is a sum.

4.2 Achievable rate, not nominal capacity

For each (attribute, rung) pair a sender encodes a uniformly random payload of k bits into that attribute using only substrate primitives, and a receiver decodes it through `list_dir`, `stat` or `read_file`. The pair is credited with k bits only if the receiver recovers the payload *exactly on every one of 8 independent trials*; k is then increased by binary search until that fails.

This is the loop shown at the right of Figure 1: encode, decode, verify, widen. It is a rate established by exhibiting a working code. It is a lower bound by construction — a better encoder would find more — and it is zero when the channel does not work, which a declared field width can never be. Where the search terminates at its ceiling the value is reported as censored ($\geq$) rather than as an estimate.

4.3 Distinguishing a closed channel from a broken encoder

A closed channel and a defective encoder produce the same output: zero bits. The measurement cannot distinguish them, so a zero is not evidence of closure on its own. Every zero reported here was therefore verified against the substrate by inspecting the encoder and the resulting events directly, and the requirement is recorded in the measurement code. This is the principal soundness risk in the method: an encoder that fails for an incidental reason is indistinguishable, in the output, from a control that works.

4.4 Unit of independence

Two agents share one substrate within a generation, so episodes within a generation are not independent — the hypothesis is precisely that one agent changes the other’s environment. The generation is therefore the unit, and all inference is computed over generations.

Because the resulting sample sizes are small, tests are exact rather than asymptotic [1]. The consequence is not cosmetic. On an outcome of 7 successes against 3 in 10 generations per arm, counting episodes gives $n = 20$ per arm and $p = 0.0256$; counting generations gives $p = 0.1789$. The pseudo-replicated analysis crosses the conventional threshold and the correct one does not.

5 Results

5.1 Measured capacity per attribute and rung

Table 1 reports achievable bits per event for every attribute at every rung.

What an uncensored cell does and does not establish. Removing a search ceiling converts a floor into a terminating measurement, and it is tempting to then call the result the channel’s capacity. It is not, and the distinction is the same one §4.3 draws for zeros.

`content` terminates because the coder writes one *character* per bit into a body limited to 8192 bytes. A body of that size holds 65,536 bits if the coder packs them, so the measured figure is eight times below what the substrate offers and what it records is the coder’s inefficiency. The `filename` and `dirname` rows are external in one factor and ours in the other: the length limit is

Attribute	open	content	filename	dirname	existence	sealed
content	8192	0	0	0	0	0
filename	1292	1292	0	0	0	0
dirname	1318	1318	1318	0	0	0
existence	471	471	471	0	0	0
size	13	13	13	13	0	0
mtime	39	39	0	0	0	0
order	0	0	0	0	0	0
count	≥ 128	≥ 128	0	0	0	0
total	≥ 11453	≥ 3261	1802	13	0	0

Table 1: Measured achievable rate, bits per event, zero-error over 8 trials. $\geq$ marks a value censored at the search ceiling. Closure is cumulative left to right. Every cell is a property of *this coder* on this substrate, not of the attribute: see the caveats below, which apply to the uncensored cells as much as to the censored ones.

the filesystem’s 255-byte `NAME_MAX`, but the per-character entropy is fixed by an alphabet we chose to be conservative. Over arbitrary bytes those rows would be roughly 2,040 bits rather than what the table reports.

We keep the conservative coders and report the smaller numbers, because L1 already commits this paper to lower bounds throughout and a bound that moves with the analyst’s ambition is not the quantity a defender needs. Censoring at a ceiling and weakness in an encoder are the same limitation reached from two directions, and neither is repaired by searching harder.

Two entries deserve comment. `order` carries 0 bits even with the channel fully open, and this is measured rather than assumed: the order coder wrote its entries in a shuffled sequence and searched for the largest permutation a receiver could recover. The receiver recovered none — `list_dir` returned the entries sorted rather than in creation order, so the sender’s sequence never reached it. The real-substrate probe (§5.10) finds the same: the stores measured do not preserve creation order in a listing. A nominal budget computed from the obvious formula, $\log_2(n!)$ at $n = 256$, would assign it 1684 bits. The gap between 1684 and 0 is the clearest available illustration of why declared widths are not capacity.

5.2 Per-attribute sums after closure

The column below adds cells measured one attribute at a time. It is an upper bound, not the residual of §4.1: it double-counts dependent attributes and inherits the censoring of any cell that stopped at a search ceiling. The residual is reported in §5.5.

After closing	Per-attribute sum
nothing (open)	≥ 11453 bits
content	≥ 3261 bits
content, filename	1802 bits
content, filename, dirname	13 bits

The benchmark task’s passcode is 20.7 bits — the dashed reference line in Figure 1. Content closure — the control a payload-inspecting monitor implements — leaves 3075 bits reachable to a single encoder (§5.5), approximately 148.7 times what an inheriting agent needs. Closing the directory names removes the two name carriers and, with them, the existence carrier — existence coding needs sender-named slots to make present or absent, and canonical names leave none. It does *not* reach zero. An object’s size is fixed at write time and read back from a listing, so it needs no sender-chosen name: an encoder reading by enumeration still recovers 13 bits after

directory-name closure, the length of the object it wrote. A coder that instead locates its object by name misses this residual once names are canonicalised and reports the rung as closed — a coder artifact, not a closed channel, and exactly the confound §4.3 names; the size figure here is measured by the enumeration coder, which does not depend on a name. The artifact exhibits both readings side by side: run at this rung, the joint encoder of §5.5 — which addresses its carriers by sender-chosen name — fails at every width and returns zero, while the enumeration coder returns 13. The disagreement is recorded rather than resolved by preference, in the `note` field of the `dirname` row of `results/capacity/joint.json`, because the two numbers measure different things and only one of them measures the channel. Reaching zero requires closing the existence carrier itself — a fixed, content-addressed manifest that refuses sender-created slots, so a sender can no longer place an object whose length it chooses — at which rung the residual is 0 bits. That zero is empirical, since no finite set of failing encoders proves non-interference (§4.3), but it is the first rung at which the encoders that reach every earlier rung recover nothing. §5.6 measures what that costs a *control* to enforce and finds it the cheaper regime; it is nonetheless the more invasive change to the cache’s *interface*, because every consumer that constructs a path by name has to migrate. Those are different costs, and §5.6 measures only the first.

5.3 Why a payload monitor cannot see a name-carried deposit

The argument, then the measurement. A monitor whose decision function reads only the bytes written cannot separate two events that differ only outside those bytes. A name-carried deposit — the payload in the entry name, the body an innocuous build artefact — is byte-identical, in the payload, to ordinary cache use. Its true-positive rate against that channel is therefore zero at every threshold, and no amount of tuning changes it. The conclusion follows from the monitor’s definition; the experiment below does not establish it but confirms that a deployed content monitor behaves as the definition predicts.

We ran that content monitor — given its best case, the full plaintext of every write — over 200 episodes in each of three classes: a content-carried deposit (the payload in the body, as before 6 July 2026), a name-carried deposit (the payload in the entry name, the body a benign artefact), and ordinary cache use. The benign bodies of the last two are independent draws from one pool of realistic build output, so any separation the monitor finds between them is separation it invents. As a positive control the monitor works: it ranks content-carried deposits above ordinary use at AUC 1.00 and catches every one at a 1.00 true-positive rate for a 5% false-positive rate. Against the name channel it is at chance: AUC 0.52 (95% interval [0.47, 0.56]), a true-positive rate of 0.00 at the same false-positive rate, and an alarm rate on name-carried deposits (0.39) indistinguishable from its alarm rate on ordinary use (0.35) — the alarms it raises are its own benign false positives, not detections. The argument said the operating characteristic must be flat; the measurement is flat.

Two further limits. The argument bounds monitors that read *payloads only*; a control that also inspects paths is outside it, and many deployed data-loss-prevention systems do inspect paths. That case can be measured rather than set aside: re-running the same joint encoder with the two sender-chosen name carriers removed — what a path-inspecting monitor neutralises — leaves 512 bits (24.8 times the passcode), carried by existence, size, mtime, still more than a 256-bit key. The existence carrier survives a path monitor because it encodes in the presence pattern of protocol-fixed slot names rather than in a chosen string; a monitor that additionally models directory *presence* drives the residual to 52 bits, an object’s size and timestamp — which are what closure removes only at the existence rung (§5.6). And the argument says nothing about how often agents would choose such a channel — §5.12 measures that separately, and finds they did not.

5.4 Which monitor sees it, and why redundancy does not

A monitor detects only what it observes, and that bounds what stacking monitors can buy. A permission guard — the class that compiles an access-control policy into enforcement — sees only whether each operation is allowed; every operation the channel uses is, so it never fires, and neither does any number of copies of it nor a debate among them, because they share one observable and so one blind spot. A content monitor is blind to a name-carried deposit for the reason §5.3 measures, and again duplication does not help. Catching such a deposit needs a monitor of a *different* observable — for instance the flow’s consequence, a cross-generation read followed by a solve too cheap to have been unaided. A monitor keyed on that consequence does separate a deposit *constructed* to be cheap from a legitimate hit, but the separation is carried entirely by that constructed difference in solve cost, and the same feature false-alarms on any legitimate cache that genuinely shortens the work — which is what a cache is for. The axis that decides detection is therefore the monitor’s observable, not the number or reasoning depth of monitors; and on this substrate the observable that would catch a deposit is the one a useful cache also trips. That points back to structural closure, not monitoring, as the containment control.

5.5 The residual, measured jointly

Table 1’s **total** row is a *sum* of cells measured one attribute at a time. That is the object this paper’s own method argues against: a total assembled from separately measured widths has never been required to carry anything, and the criterion of §4 — credit k bits only when a receiver recovers k bits — was never applied to it.

Worse, two of the summed cells are not independent. The existence coder writes one file per set bit, so the number of entries is a deterministic function of the existence pattern: the bits credited to **count** are already spent by **existence**, and adding them double-counts.

We therefore measured the residual with a single encoder. One payload is split across 5 carriers — **dirname**, **filename**, **existence**, **size** and **mtime** — written into one substrate state with content closure applied, and recovered by a receiver using only **list_dir** and **stat**. **count** is excluded as dependent, **order** because it carries nothing.

The joint encoder recovers 3075 bits, exactly on 8/8 trials. That is the figure this paper reports as the residual after content closure. It is 186 bits below the summed total, and that gap decomposes into two parts, both measured:

Overstatement by summing per-attribute cells	186 bits
the count dependency, excluded by construction	128 bits
joint encoder loss, measured by search	58 bits

The first term is the dependency predicted above: an existence pattern determines the entry count, so the joint coder excludes **count** rather than double-counting it. The second is not predicted by anything. It is what one encoder loses relative to five separate ones when it has to write all of them into a single substrate state, and it is reported as loss rather than explained.

Neither term is asserted. The joint coder reads each carrier’s width from the per-attribute measurement rather than holding it as a constant, so the two figures are computed from one artifact and the gap between them is whatever the search finds. A decomposition of this kind is only meaningful if it cannot be tuned: fixing the widths independently of the measurement would allow the gap to be set rather than observed.

At the **filename** rung the same comparison gives a loss of 0 bits: with only **dirname** and **existence** live, one encoder carries everything the two separate measurements do. The carriers interact at one rung and not at the other, which is a property of the substrate rather than of the method.

3075 bits is 148.7 times the 20.7-bit passcode. The qualitative conclusion is unchanged and now rests on a number that meets the criterion: closing content leaves an order of magnitude more capacity than an inherited answer requires.

5.6 What each rung costs to enforce, and who can check it

A residual with no cost beside it is a recommendation a lab can decline for free. This section supplies the cost, and is explicit about what kind of cost it is.

It is not engineer-weeks. What a production integration would cost a team in engineer-weeks is not estimated here; inventing that figure would be exactly the unmeasured quantity §5.11 argues against. What can be shown is that the control runs on a real store. We deployed one of the recommended controls — name canonicalisation, content-addressing the write path, the discipline an OCI registry already applies to its blobs — as a write-path shim on a live commercial object store, at 1 enforcement point, and confirmed both that the name channel is real there and that the shim discards the sender’s chosen key while a receiver that knows the content it wants still fetches it by hash. This is a deployment demonstration, not a sampled efficacy rate: once the key is a content hash the sender cannot choose it, so the name channel closes by construction. Its scope is that construction. An S3 key is *flat*, so one canonicalisation closes the filename and directory-name carriers together; a store with hierarchical, sender-chosen names does not get this for free. The OCI registry of §5.10 content-addresses its blobs and still carries 331 bits in sender-chosen *repository* names, so the control must canonicalise every sender-choosable segment, not only the leaf. The human integration effort remains the one quantity this section does not measure; two more it does, and an auditor who never saw the lab’s network can check them, are these.

Mediation points — how many distinct substrate primitives a control must sit on. This is the integration surface, it is fixed by the API rather than by the implementation, and an auditor can count it from a specification.

Enforcement class — what the control must *do* at those points. A **structural** control enforces a naming or namespace invariant at the API boundary: constant work per call, no artefact body is read, and there is no operating point to tune, so it has no false positives. A **payload** control reads or rewrites bodies: work scales with bytes rather than calls, and it needs a threshold, which is where false positives and false negatives enter. These are different cost regimes, not two points on one scale.

Table 2 gives both axes for every rung, beside the residual each one leaves. Read down the mediation column for what a control costs to integrate, and across the payload-class column for which regime it sits in; the two together are what §5.6 means by cost.

Rung	Rows closed	Mediation points	Payload-class rows	Bits
content	1	2	1	3261
filename	2	5	1	1802
dirname	3	6	1	13
existence	4	7	1	0
sealed	8	7	2	0

Table 2: Cost of *holding* each rung: every row at or before it stays closed simultaneously. Mediation points are de-duplicated across rows.

Closure is not ordered by cost, and the incident shows it. The response in July was not a closure at all. OpenAI records that its security team “blocked a known privilege-escalation route, removed exposed credentials, and later rebuilt Artifactory,” and that clearing the message board was “the side effect” of the rebuild [17]. That is credential revocation and reinstantiation

— in this paper’s terms a *wipe*, which erases state without removing any attribute’s capacity to carry the next message. Every row of Table 1 was as open after the rebuild as before it.

The distinction matters for cost because a wipe is recurring and a closure is not: the operator pays again at every remediation, and gets a substrate with the same residual each time. Content closure — the control a payload-inspecting monitor implements, and the one an operator reaching for a durable fix would most likely choose next — is a **payload** control in the expensive regime, 2 mediation points, and would leave 3075 bits reachable to a single encoder. Directory-name closure reduces this to 13 bits, an object’s size; reaching zero requires existence closure, a **capacity** control — a fixed content-addressed manifest that refuses sender-created slots — holding 7 points in total, of which 1 is payload-class, the one the lab has already paid for. The distance from the content control is **5 additional mediation points**: no further artefact body is read, no threshold is chosen, and nothing scales with stored bytes. The rung that reaches zero is still in the cheaper regime than the content control that does not.

This inverts the usual adoption argument. The obstacle is not that the effective control is expensive; it is that the ineffective one is the one that looks like security.

Verifiable without the lab’s network, and cheap enough to adopt. Two properties decide whether a control of this kind is worth writing down: an outside party must be able to check compliance without privileged access, and a lab must be willing to pay for it. Both are answerable here from artefacts rather than assurances. On the first, the ladder is a script and a substrate: it runs with no model, no network and no key, so an auditor reproduces the residual from a published configuration on their own hardware and compares it against the lab’s declared rung. What is checked is a property of the specification the lab published, established by exhibiting a working code against it — not an attestation the lab makes about itself. On the second, §5.7 prices every rung in the currency a lab actually spends: the closure that reaches zero costs **1.00** task success, and costs **0.00** more than directory-name closure, which already leaves only 13 bits. The evidence for each clause is listed in Appendix B, and in every case it is an artefact the lab already holds.

The scope of that claim needs stating precisely. The July response was a wipe, so the paragraph above is a counterfactual about the control an operator would reach for next, not an observation of one a lab chose. What the record supports is narrower: the response taken was recurring rather than durable, and the durable control most readily reached for is the one in the expensive regime. Whether a lab would in fact choose content inspection over namespace separation is a prediction this paper does not test.

A third party can verify this without the lab’s network. The criterion an evaluation-infrastructure control has to meet is that someone outside can check it. Both quantities above are external: the mediation count comes from the substrate’s published API, and the enforcement class from what the control is specified to do. Neither requires access to traffic, logs or hosts.

The residual itself is externally checkable on the same terms. A lab publishes its substrate configuration and its closure disposition per attribute; an auditor re-runs the released encoder against that configuration and either recovers a payload or does not. What the auditor verifies is a property of the *specification the lab published*, established by exhibiting a working code against it — not an attestation, and not a number the lab computed about itself. A lab whose published configuration does not match what it runs has misdeclared its configuration, which is a different and more visible failure than an unverifiable safety claim.

5.7 What each rung costs the agents that use the cache

Mediation points price the *defender’s* work. They say nothing about the price paid by the agents the cache exists to serve, and that is the price a lab weighs when deciding whether to adopt a

control. A separate arm supplies it: 5 seeds, 40 episodes per cell, probe budgets 3–16, a single agent with an oracle cache user, no channel and no inheritance. Table 3 reports task success as each rung is added.

closure	probe budget													
	3	4	5	6	7	8	9	10	11	12	13	14	15	16
nothing closed	1.00	1.00	1.00	1.00	1.00	1.00	1.00	1.00	1.00	1.00	1.00	1.00	1.00	1.00
content	0.00	0.00	0.00	0.00	0.00	0.00	0.03	0.07	0.15	0.28	0.42	0.62	0.90	1.00
+ filename	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.03	0.07	0.15	0.28	0.42	0.62	0.90
+ dirname	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.03	0.07	0.15	0.28	0.42	0.62	0.90
+ existence	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.03	0.07	0.15	0.28	0.42	0.62	0.90

Table 3: Task success under cumulative closure. Rows below **+ filename** are identical to it: once names are closed, every further rung is free.

The closure that works is the one that costs nothing. Closing **dirname** in isolation costs **0.00** task success at every budget swept, and so do **dirname**, **size**, **order**. Cumulatively the point is sharper still: adding **dirname** closure on top of **filename** closure changes task success by **0.00** at every budget. The step that takes the residual from 1802 bits to 13 is free.

The expensive closures are the ones that do not finish the job. Content closure costs up to **1.00** task success and leaves ≥ 3261 bits reachable; filename closure costs up to **1.00** and still leaves 1802. Both are paid in full before any of the capacity that matters is removed. A lab that prices its controls by intuition would adopt the cheap-looking one and pay the most for the least; the ordering here is the reverse of the intuitive one, and it is measured rather than argued.

Two limits on this table. Above a budget of about 14 an agent can simply recompute what it needed, so no closure can show a cost there and the right-hand columns understate every control. And the cache user is an *oracle*: it reads what a real model might fail to find. Success under closure is therefore over-estimated and every cost above is a **lower bound** on what a real model would pay. We report it that way rather than presenting a floor as an estimate.

5.8 What the residual is worth, measured against secrets we did not define

Comparing the residual to this study’s own benchmark passcode has a defect: the substrate’s parameters and the passcode are both ours, so a multiple built from them is partly circular. Table 4 anchors the comparison instead to payload sizes fixed by third parties, and asks of each whether the jointly measured residual is large enough to carry it.

Secret	Bits	Fits in 3075?
IPv4 address and port	48	✓
AWS access key ID	80	✓
UUID version 4	122	✓
GitHub personal access token	214	✓
Ed25519 public key	256	✓

Table 4: Information content of identifiers and credentials whose sizes are set by external specification, against the jointly measured residual after content closure. Sizes are information content, not encoded length: a UUIDv4 is 36 printable characters but only 122 bits vary.

5 of 5 fit, which is to say the table does not discriminate. At 3075 bits every secret listed fits in a single deposit, including the largest.

That is the result, and the table is retained for it. The question it answers — whether the residual is large enough to matter — is settled more decisively by every row fitting than by some subset of them. What the table cannot do is rank closures by which secrets they stop, and it should not be read that way.

The comparison is also conservative in a way worth stating. Nothing in the threat model limits a sender to one deposit, and the incident involved over 70,000 messages [12]. A residual that carries the largest of these in a single substrate state carries anything at all across a handful of them, so the single-shot framing understates the channel and is used for that reason.

We are not claiming these particular secrets moved in July. The public record says agents exchanged credentials, commands and results without enumerating formats. The sizes are useful only because nobody involved in this study chose them.

5.9 The control this ladder is measured against

The cheapest objection to everything above is also the strongest, and it should be stated in its most damaging form: if giving each agent its own namespace converts co-tenant flow into an ordinary access attempt across a boundary where authorization controls already work, then the ladder instruments a problem a one-line fix dissolves, and a lab should adopt the fix rather than the instrument.

On the capacity axis the objection is correct, and we measured it rather than argued with it. Namespacing is not a further rung. The ladder closes *attributes* on a shared resource; namespacing removes the sharing, so it acts on a different axis and is strictly stronger on the one it acts on. Across all 8 carriers, with a sender encoding into one namespace and a receiver enumerating its own, 0 of 1153 planted artefacts were observable from the receiver’s side and 0 of 320 escape attempts resolved — against a same-namespace positive control of 11,453 bits, so the apparatus was measuring confinement rather than its own inability to address anything.

Two things that failure would have looked like are worth naming, because both occurred during construction. A first cross-namespace measurement handed the coder the receiver’s substrate, so it wrote and read in one namespace and reported a spurious leak on seven carriers; the number it produced was the top of its own search range. And a first confinement, implemented by prefixing the path and deferring to the ordinary containment check, was escaped 320 times out of 320: `../agent_a` resolves to a sibling subtree that is still inside the root, so the check passed. Containment has to be against the namespace, not the root. The probe attempts escapes for exactly this reason rather than assuming the control holds.

What it costs is reuse, not integration. On the two axes §5.6 applies to the ladder, namespacing is cheap: 1 mediation point, because every primitive already routes through path resolution, and the structural enforcement class, which is $O(1)$ per call with no operating point and therefore no false-positive surface. The cost lands somewhere else. A shared cache exists so that two agents needing one artefact fetch it once; partition it per agent and that property is gone by construction. Table 5 counts fetches over 200 workload draws per row:

The cost is negligible when agents share little ($1.11\times$ at two agents) and severe when they share a lot: at 16 agents the same workload goes from 100 fetches to 640, a factor of 6.40. Overlap is not incidental to the deployment; it is the reason the cache is there.

So the recommendation this paper makes is not the ladder over namespacing. It is namespacing where the substrate can be partitioned, priced as above, and the ladder for the residual — shared build artefacts, content-addressed blobs, any store whose value is precisely that two agents reach the same object. The incident’s substrate was of that kind. What the ladder supplies there is a measurement of what a partial closure buys, in bits, so that accepting a residual is a decision rather than an oversight.

agents	pool	each	overlap	shared	per-agent	×
2	100	20	0.10	36	40	1.11
4	100	20	0.26	59	80	1.36
8	100	20	0.48	83	160	1.93
16	100	20	0.70	97	320	3.29
16	100	40	0.84	100	640	6.40
16	50	20	0.84	50	320	6.40

Table 5: Artefact fetches under one shared cache against per-agent namespaces. The multiplier rises with overlap — the same property that motivates having a cache at all.

5.10 The same questions asked of real infrastructure

Everything above is measured on a substrate we wrote, and L3 concedes that its parameters are ours. The obvious objection is that deriving a search ceiling from a constant we chose relabels the circularity rather than removing it. This section answers it by asking the same questions of a real shared volume and a real artifact store, both in containers, with the limits *discovered* rather than declared.

A real shared volume. Two containers with separate PID and mount namespaces, sharing exactly one named Docker volume — the isolation an evaluation harness claims to provide, and the one thing it does not isolate. Probing that mount through ordinary file operations:

Quantity	Synthetic (declared)	Real (measured)
Longest path component	255	255
Largest body accepted	8,192	67,108,864
Entries in one directory	256	200,000
Timestamp resolution	whole seconds	1 ns, exact
Listing preserves creation order	no	no

Two of these overturn a figure reported earlier. The body limit is 8,192 times what the synthetic substrate declares, so `content` is not the 8192-bit channel Table 1 reports — that number is our constant, and the coder’s inefficiency on top of it. And a real mount round-trips timestamps at 1 ns, exact, where the synthetic substrate exposes whole seconds: roughly thirty additional bits per event on a carrier no payload monitor inspects.

Two are confirmed rather than overturned. `NAME_MAX` really is 255, and a real listing really does not preserve creation order — so the `order` row’s measured zero is a property of filesystems and not an artefact of our model.

A real artifact store. An OCI distribution registry, default configuration, authentication disabled — the configuration a single-tenant internal harness typically runs. Table 6 gives the per-attribute verdict; each one was established by performing the operation, not by reading documentation.

The registry closes `filename` properly and by construction: a blob’s identifier *is* its `sha256`, so a sender cannot choose it. That is filename closure shipped as a product decision rather than as a security control, and it is the only rung the store closes.

Because the ladder is cumulative, closing a middle rung earns nothing. The store’s native rung is **open**. A sender chooses the repository name — 331 bits over the grammar the distribution spec permits — and every co-tenant reads it back from `GET /v2/_catalog`. Existence carries too: 16 of 16 pre-agreed repository names were recovered exactly from the catalog, with no payload written at all.

Attribute	Verdict	Evidence
<code>content</code>	carries	512-byte blob round-tripped: True
<code>filename</code>	closed	blob identifier is its own sha256 (sha256:40c1fcca7ab33fdf...); the store derives it from content and rejects a mismatch, so it is content-addressed by default
<code>dirname</code>	carries	pushed repository ujzde8gxd6ncf10e... and found it in /v2/_catalog
<code>existence</code>	carries	16/16 agreed slots recovered exactly from the catalog
<code>size</code>	carries	Content-Length header present: True
<code>mtime</code>	closed	Last-Modified header present: False

Table 6: What a real OCI registry carries between co-tenants, probed directly. Generated from `results/real-store-registry.json`.

This is the same shape as the channel the July agents used once content was wiped — names of containers rather than contents of files — and it requires no misconfiguration to exist.

An object store removes two carriers and widens a third. The same probe run against a commercial object store (eu-north-1) returns a different substrate, and the differences are structural rather than a matter of degree. A sender cannot set the timestamp a reader observes: the store assigns it and exposes no API to override it, so `mtime` is not a small carrier here but an absent one. Nor is `dirname` separable from `filename`, because a directory is a prefix of the key rather than an object — so the two rungs the ladder treats as successive are a single control. Ordering is unobservable for the same reason it is on a mount — keys created `zz`, `mm`, `aa` are returned `aa`, `mm`, `zz`.

Against that, the name carrier is far *wider*: the store accepted a 1,018-character key segment, 4.0 times the 255 a path component allows, with no directory-entry ceiling to bound it — 1,200 entries were written under one prefix and all 1,200 were listed back. Bodies reached 8,388,608 bytes at our probe ceiling rather than at the store’s.

The ladder’s shape is a property of the substrate, not of the method. This is the finding we did not anticipate and it qualifies the headline. On a filesystem, closing `filename` leaves 1802 bits, `dirname` leaves 13, and only `existence` closure — a fixed content-addressed manifest — reaches 0; on an object store those are not decisions a defender can take in sequence. A lab whose harness runs on a mount faces the four-rung ladder reported here. A lab whose harness runs on an object store faces a shorter one, reaches zero with one control rather than two, and pays less than §5.7 implies — but has a wider name carrier to close while the ladder still stands. Neither substrate is the one the ladder was measured on, and a residual quoted from Table 1 transfers to neither without re-measurement.

What this does and does not establish. It does not re-measure the ladder against real infrastructure; the capacity figures in Table 1 remain synthetic, and §5.6’s costs remain structural. What it establishes is narrower and, we think, more useful: the attributes the ladder enumerates are present, sender-chosen and receiver-observable on infrastructure nobody in this study configured; the store most likely to be found under an evaluation harness sits at no rung of the ladder at all; and on an object store two of those attributes are not carriers in the first place, which is a statement about the substrate that no measurement on a mount would have produced.

5.11 Why the fraction removed is not reported

It is tempting to report that content closure removes most of the capacity. That figure is not stable, and this project has watched it move. Raising the per-attribute search ceilings from hand-chosen values to ones derived from the substrate’s own constants changed the open-rung total from ≥ 4297 bits to ≥ 11453 and the content-rung sum from ≥ 201 to ≥ 3261 — while the substrate itself was not touched. A percentage built on either pair would have moved with it, and a reader could not tell the movement of the analyst from the movement of the channel.

That movement can be reproduced on demand rather than left as history. Holding the substrate fixed and sweeping only the analyst’s search ceiling on the censored content cell (Table 7), the fraction content closure appears to remove slides from 0.24 to 0.71 — a factor of three — as the ceiling rises, and plateaus only once it passes the substrate’s own limit. The joint residual, measured by the same encoder and never touching the content cell, holds at 3075 bits throughout. A percentage reported from any single ceiling is a reading of the analyst’s knob; the residual is a reading of the channel.

Content ceiling (bits)	Content cell (bits)	Open total (bits)	Fraction removed	Joint residual (bits)
1024	1024	4285	0.24	3075
2048	2048	5309	0.39	3075
4096	4096	7357	0.56	3075
8192	8192	11453	0.71	3075
16384	8192	11453	0.71	3075

Table 7: The fraction is the analyst’s; the residual is the channel’s. The substrate is held fixed and only the content-cell search ceiling varies; above the substrate’s own limit the ceiling stops mattering and the open total plateaus, while the joint residual never moves.

The figure that is immune is the joint residual, and only because it excludes the censored cell. 3075 bits were recovered by one encoder across carriers whose measurements *all terminated* — `dirname`, `filename`, `existence`, `size` and `mtime`. The one row still censored at its ceiling is `count` (≥ 128), and the joint coder omits it by construction because an existence pattern determines it. So the residual this paper reports does not contain a censored quantity, and raising a ceiling cannot move it.

The per-attribute *sums* in Table 1 do contain that censored cell, which is one more reason not to quote them as the residual. They are reported because the ladder needs a per-attribute view, not because any encoder was asked to carry them.

5.12 The behavioural arm: agents read the substrate and never write to it

Everything above concerns what a substrate *can* carry. Whether agents in fact use it is a separate question, and the behavioural arm answers it in a way the design did not anticipate.

Capability. A model that cannot perform the task at all produces the same output as a model that declines to use the channel: no deposit, and a null Δ . The two are indistinguishable after the fact, exactly as a broken encoder is indistinguishable from a closed channel in §4.3. Arm membership is therefore gated before any episode is run, by two probes that isolate the capacities an episode presupposes. The first requires the model to emit a well-formed tool call, since an agent that cannot call `validate` cannot act on the substrate whatever its intentions. The second requires the elimination inference the search task is built on, sampled more than once because a single pass is not evidence: a model that passed this probe on one attempt and failed it twice thereafter would otherwise be admitted on the strength of the first. A model failing either probe is excluded. The gate bounds one reading of the null reported below — that it came from agents

incapable of depositing — without establishing that the agents which remained would have deposited under any condition.

Calibration. The inheritance advantage Δ is only interpretable if the baseline is neither pinned at the floor nor at the ceiling: at the ceiling every agent succeeds unaided, and at the floor no generation-1 agent solves anything, so nothing is deposited and generation 2 inherits nothing. Either way Δ is zero by construction. The pre-registered gate therefore requires solo success strictly between 0 and 1 in the `no_substrate` condition, which contains no shared store and so cannot express the effect it is used to license.

One model has been excluded on the floor, and the exclusion is a finding rather than a discard. Qwen/Qwen3-8B solved 0 of 3 baseline episodes at probe budget 12, an upper 95% Wilson bound of 0.56. It passed the capability probes while doing so: 0 malformed probes and 0 invalid tool calls across those episodes, with each submission reasoning explicitly from validator feedback about which positions were fixed. That separation is the point of running two gates rather than one. The floor is a statement about the difficulty of this task family for that model, not about its ability to act on a substrate, and a single gate could not have told them apart. Replacing it with a frontier-class model moves the anticipated risk to the opposite end — optimal play resolves a four-position, ten-symbol code in roughly five to six probes — which is why the budget below was recalibrated downward rather than upward. That ceiling is a prediction the design guards against, not an exclusion this study observed.

The selection screen. 3 models were screened as candidates for the reported calibration model and 1 is reported. Table 8 lists all of them, because a survivor presented without its screen reads as a model that was used rather than one that was chosen.

Model	Route	Solo	Verdict	Status
qwen8-hf	local (transformers)	0/3	floor	screened out
claude-haiku-4.5	OpenRouter	10/31	gate passed	reported
gemini-2.5	Gemini Developer API	1/1	undetermined	screened out

Table 8: Models screened as candidates for the reported calibration model. *Solo* counts only episodes that completed without a provider error, since a run cut short by a rate limit is not evidence about a model.

The gate resolves per operating point, not per model. A single rate for a model is the wrong summary, because the gate is a property of the operating point: a model can floor at one probe budget and sit inside the window two budgets later, and pooling the two would report a rate strictly inside $(0, 1)$ while no single operating point is, manufacturing a gate pass out of results that individually fail it. Table 9 makes the difficulty curve explicit, sweeping the probe budget for the reported model and for a set of further models run on Bedrock — whose per-token billing, unlike a shared free tier, lets the budget be swept in full without a rate limit truncating the run — across 6 model–route rows over 7 probe budgets, with the operating points satisfying the gate in bold.

Read across a row and the difficulty curve is visible where a pooled rate would hide it. `claude-haiku-4.5` (OpenRouter) at 20, 24, 28; `gpt-oss-120b` (Bedrock) at 16, 20, 24; `kimi-k2.5` (Bedrock) at 12, 16, 20, 24. The reported model is not the only one that clears the gate — two open-weight models from different families do so, each across at least three consecutive budgets. These two rows were swept to nine seeds rather than three, for the reason the calibration of §5.12 gives: a window located at $n = 3$ can be sampling noise, so a gate pass is reported only where it survives the larger sample. The screen therefore selects among models that pass rather than identifying the only one that can.

model	route	probe budget						
		4	8	12	16	20	24	28
claude-haiku-4.5	OpenRouter	–	0/3	0/3	0/3	1/6	4/8	5/8
gpt-oss-120b	Bedrock	0/9	0/9	0/9	6/9	5/9	5/9	–
amazon-nova-lite	Bedrock	0/2	–	0/2	–	–	0/2	–
amazon-nova-pro	Bedrock	0/3	0/3	0/3	0/3	0/2	0/3	–
kimi-k2.5	Bedrock	0/9	0/9	4/9	6/9	5/9	8/9	–
qwen3-32b	Bedrock	0/3	0/3	0/3	0/3	0/3	0/3	–

Table 9: Solo success in `no_substrate`, solved over clean episodes. API-error episodes are excluded rather than counted as failures. Bold marks an operating point strictly inside $(0, 1)$, which is the gate condition.

The 2 exclusions are not of one kind, and the difference matters for what the screen licenses. One is a property of the model: it solved nothing at the budget swept, having first demonstrated it could act. The other is a property of the *infrastructure* — a free-tier request cap reached before enough clean episodes accumulated to place the model inside or outside the window. It is recorded as undetermined rather than as a failure, because nothing was measured about it: a provider withdrawing capacity says nothing about whether an agent would deposit. A model that a provider limit leaves undetermined on a throttled endpoint can still be measured where the budget can be swept in full, which is what Bedrock’s per-token billing provides and why the grid above extends the screen rather than repeating it.

The screen therefore narrows the field on a criterion measured where the effect cannot appear, but it does so unevenly, and the reported model is the one that both cleared the window and had a provider willing to serve enough episodes to show it.

Table 10 reports 31 baseline episodes for anthropic/claude-haiku-4.5 across 6 probe budgets.

Probe budget	Solved	Rate	95% Wilson CI
8	0 / 3	0.00	(0.00, 0.56)
12	0 / 3	0.00	(0.00, 0.56)
16	0 / 3	0.00	(0.00, 0.56)
20	1 / 6	0.17	(0.03, 0.56)
24	4 / 8	0.50	(0.22, 0.78)
28	5 / 8	0.62	(0.31, 0.86)

Table 10: Solo success in the `no_substrate` baseline. Rates strictly between 0 and 1 are in bold.

The operating point is the budget nearest the pre-registered target of 0.45: **24** probes, at solo success **0.50** over 8 episodes, 95% Wilson interval (0.22, 0.78).

Three seeds per budget were not enough to locate this. At $n = 3$ the sweep was not monotone — budget 24 returned 3/3 and budget 28 returned 1/3 — and a larger budget cannot make a task genuinely harder, so that ordering was sampling noise, part of it task variance since the seed selects which code the agent must find. Extending to eight seeds resolved it into the monotone curve above. A difficulty sweep at $n = 3$ measures noise in the shape of a result.

The matrix. 40 episodes, `open` and `wipe`, 10 generations $\times$ 2 agents per arm, at the calibrated budget, with 0 API errors.

	open	wipe
Episodes	20	20
Read the substrate	38 / 40	
Wrote to the substrate	0	0
Task success	11 / 20	13 / 20
Generations with a success	8 / 10	8 / 10

The deposit rate is zero, and that is the result. 0 of 40 episodes wrote to the substrate; the at the generation unit that is 0 of 20 cells, upper 95% Wilson bound 0.161. (Counted over episodes the bound is 0.088, but episodes within a generation share a substrate and §4.4 rejects that unit; the looser figure is the honest one.)

The design did not fail for want of power. It failed because the condition hypothesised to drive deposits was never presented: no unsolvable task was drawn (see below). A null measured with the treatment absent is not a null about the treatment, and neither bound should be read as one.

The consequence is structural. The inheritance advantage is the last term of a conjunction,

$$\text{deposit} \rightarrow \text{discovery} \rightarrow \text{use} \rightarrow \text{success},$$

and only the last term was measured. With no deposit in any episode, the substrate an inheriting agent encounters holds nothing but what the harness seeded — so **open** and **wipe** are not two conditions. They are the same condition run twice, and their difference estimates the difference between a distribution and itself. Generation-level success was 8/10 in both arms, Fisher exact $p = 1.0000$, which is what that identity predicts. **Δ is undefined for this configuration rather than small**, and no sample size repairs it.

What the agents did instead. 38 of 40 episodes inspected the shared store, at a median of 5 read calls each, and none wrote to it. Agents treat the cache as a resource to read, not as a channel to use. The asymmetry is the more informative half of the run: it is a behavioural claim the capacity measurement cannot produce, and it bears directly on containment, since a control that prevents writing addresses a behaviour these agents did not exhibit.

Why the deposit never happened. The task assignment gives an agent a blocked — that is, unsolvable — task only when its index satisfies $i \bmod 4 = 3$. The matrix ran with 2 agents per generation, so $i \in \{0, 1\}$ and that branch is unreachable: 0 blocked tasks were drawn in 40 episodes.

This is not a bookkeeping detail. The agents in the July 2026 incident deposited under the pressure of problems they could not solve; a large fraction of that benchmark’s problems were unintentionally impossible, which is what made coordination rational. The matrix removed exactly that pressure. The pre-registered prediction that deposit rate is higher on blocked than on solvable tasks was therefore not refuted — it was untestable in this configuration. §5.13 reports the run that tests it.

5.13 The follow-up: deposit behaviour under unsolvable tasks

Raising the generation width to 4 agents makes the blocked branch of the assignment rule reachable without changing anything else: same model, same calibrated budget, same substrate, same closure. 40 episodes were attempted and 35 completed across 9 generations, of which 8 drew an unsolvable task and 27 a solvable one. 19 episodes succeeded, and — as the design requires — 0 of the blocked tasks did.

The treatment was therefore present, and the outcome did not change. 0 of 35 episodes deposited anything for a successor, giving an upper 95% Wilson bound of 0.30 at the generation unit; among the 8 blocked episodes specifically the bound is 0.32. Reading was again near-universal: 35 of 35 episodes inspected the shared store, at a median of 6 read calls. **The pre-registered prediction is now tested rather than untestable, and it is not supported at this sample size.** The bound is wide, and 8 unsolvable tasks cannot exclude a deposit rate that a larger run would detect; what the run excludes is the explanation that the earlier null was an artefact of never presenting the pressure.

A writability probe, which is not a deposit. 1 episode wrote to the shared store, in 1 of the 9 generations. Classifying it correctly matters more than its rarity suggests, because it is the entire difference between a deposit rate of zero and one of $1/35$, so the classification is made by a rule over the event log rather than by reading the transcript. A sender leaving a message for a successor has no reason to read its own write back — the write is the act. An agent establishing whether a store is writable has every reason to, because confirming persistence is the act. The rule is accordingly: a write followed immediately by the same agent reading or stat-ing the same path is a probe.

The single write meets it. An agent holding an unsolvable task, after twenty turns of failed validation, wrote four bytes to the cache, read them back, and stat-ed the result in three consecutive calls, then continued searching and submitted nothing. The content was a placeholder, not task knowledge, and no later generation referenced the path.

The episode is nonetheless the most informative single observation in the behavioural arm, because it separates two failure modes that a deposit rate of zero otherwise conflates. An agent that never discovers the channel and an agent that discovers it and declines to use it are different problems with different remedies. This agent discovered it: unprompted, from task frustration, with nothing in the task description suggesting the cache was writable, it established the affordance the entire capacity measurement is about and then had no use for it. Discovery is not the barrier. What is absent in a 40-episode benchmark, and present in the incident, is a reason to send — a persistent population, a shared objective, and a successor an agent has grounds to believe is looking. That is a claim about conditions, and it is testable; $n = 1$ makes it a hypothesis for a larger run, not a finding.

That larger run exists, and it was published while this paper was being written. Paglieri et al. [18] report 100 autonomous agents pursuing a shared objective under competitive pressure, in which an exploit discovered by a single agent propagated through a shared knowledge library and then peer-to-peer across the collective. Their three conditions are exactly the three named as absent above, and their agents deposited: one records having “proudly logged and immortalized this discovery” to a file in its persistent memory “to preserve it for future iterations,” and others catalogued the technique in local reference wikis for the same purpose. That is the behaviour this arm measures at zero, observed, with the sender’s motive stated in its own transcript.

Their study also names the mechanism, which matters more than the population size. Agents defected under an *asymmetric lockout*: a problem was permanently closed to the first agent that solved it, so as the pool drained, agents that declined to cheat faced exclusion rather than merely a lower score. Scarcity, not scale, is what converted a stated rule into — in one agent’s words — “a bluff.” Our benchmark has no such structure: an agent that fails costs itself nothing that another agent gains.

We did not predict this and claim no credit for it; what it supplies is the other half of a bound. The finding here is not that agents decline to use shared infrastructure to coordinate — given a reason to send, they demonstrably do. It is that absent one, they do not, even when handed unsolvable tasks and a writable store they were observed to probe. The behavioural question is about conditions rather than capability, and §4.1 says what the channel carries once those conditions are met. It also makes the next experiment specific rather than merely larger:

introduce first-to-solve lockout and measure the deposit rate again.

Censoring. 5 of the 40 attempted episodes ended on an HTTP 402 when the account’s credit was exhausted, all in the final two generations. They are excluded from every rate above rather than counted as failures, since the terminating condition is external to agent behaviour; counting them would understate the success rate and inflate nothing that matters to the deposit result, which is zero on either treatment.

One qualification on that argument. Independence from the *outcome* is not the only property at stake: the losses fall in the last two generations, and §4.4 makes the generation the unit of inference, so they are not missing at random with respect to that unit. Generation 10 contributes no cell at all and generation 9 contributes a cell of three rather than four. The effect on a deposit rate of zero is negligible in any direction, but the censoring is clustered and calling it external to agent behaviour does not by itself make it ignorable.

5.14 A significant result on a comparison with no treatment in it

The matrix supports one further analysis, and it produces a positive result that we report in order to reject.

The design is matched by construction. The seed selects the task, both arms run the same seeds, and in 20 of 20 cells the two arms drew the same task. Pairing on the cell therefore removes task difficulty — much the largest source of variance in how many oracle probes a solution costs — from the comparison. On the 7 cells where both arms solved, the median cost was 16 probes under *open* and 20 under *wipe*, a median paired difference of 5 probes favouring *open*.

The differences run one way. 6 pairs favour *open*, 0 favour *wipe*, 1 is tied; an exact two-sided sign test gives $p = 0.031$. A matched-pairs design, a directional hypothesis fixed in advance, and $p < 0.05$: by the conventions that govern most of this literature, that is a result.

It is not one, and the rest of this paper is what establishes that. The hypothesised mechanism is that an agent inherits something a predecessor left. 0 of 40 episodes left anything, and the harness re-seeds after a wipe, so an agent in either arm opens a cache holding the same three empty directories. The substrate an *open* agent encounters and the substrate a *wipe* agent encounters are the same substrate. The treatment does not exist, and a difference between two draws from one distribution is what remains.

Three mundane explanations fit. 7 pairs is a small sample, and a sign test on 7 observations reaches $p < 0.05$ on a 6–0 split alone. The arms also ran sequentially rather than interleaved, roughly half an hour apart against a hosted model, so any provider-side variation over that interval is confounded with the arm. And the pairs are the very unit §4.4 rejects: 7 episode-level cells drawn from at most 10 generations are not independent, because episodes within a generation share a substrate — the same pseudo-replication this paper warns against, reproduced in its own post-hoc analysis. Any of the three suffices, and the third is the one this study is on record as knowing about. That is the point: the uninteresting explanations are the ones left standing, and nothing in the p -value distinguishes them from the interesting one.

We report this because the alternative is worse. Had the deposit rate been unmeasured — had we recorded only success and cost, as an evaluation harness naturally would — this analysis is what the study would have produced, and it would have read as a modest confirmation of the pre-registered hypothesis. The zero deposit rate is the only thing standing between that number and a publishable claim. This is the same discipline §5.11 argues for on the capacity side: report the quantity that constrains the conclusion, not the one that flatters it.

6 Threats to validity

Censoring. Five of the open-rung measurements terminate at their search ceiling and are lower bounds. Totals inherit that status and are reported with $\geq$.

Encoder strength. Every figure is a lower bound produced by one encoder. A stronger encoder raises the residual; none of the closure conclusions weaken as a result, since they assert that capacity *remains*, not that it is bounded above.

Constructed senders. The monitor comparison (§5.3) uses constructed rather than sampled behaviour. It supports a claim about observability, not about propensity.

Structural detection is not established. §5.4 argues that catching a name-carried deposit needs a monitor of a different observable than the content and permission monitors have, and that a monitor keyed on the flow’s consequence can separate a deposit *built* to be consequential from a legitimate hit. That separation is carried entirely by the constructed difference in solve cost: the same feature false-alarms on any legitimate cache that genuinely shortens the work, and because §5.12 recorded no real deposit, whether such a monitor separates real behaviour is unmeasured. No clause should cite it as evidence that structural monitoring works in deployment.

The behavioural effect is undefined, not small. Whether a shared substrate raises success for an inheriting agent is the quantity that would connect this capacity to observed behaviour, and §5.12 reports why the matrix could not estimate it: the deposit rate was 0 of 40, so the two arms were the same condition. That is a bound on this configuration, not on agents in general — one model, one prompt, one task family, two agents per generation, and no unsolvable task drawn. Nothing in §5 depends on it, because the capacity ladder is measured without any model call.

A second model was attempted, and the gate refused it. The obvious remedy for a single-model null is a second model, so we ran one: `qwen3-32b`, a different checkpoint reached through a different provider, on the identical harness and task family. It never cleared the calibration floor. Across 6 probe budgets from 8 to 28 it solved 0 of 18 baseline episodes, and the pre-registered rule — run the matrix only where solo success falls strictly inside $(0, 1)$ — stopped the run before any matrix episode was spent.

This is worth reporting rather than discarding, for three reasons. It is a clean floor and not an artefact of the provider: 0 of the 18 episodes ended in an API error, so the model was asked the question 18 times and failed it every time, at 482,694 tokens. It is a direct instance of the failure §5.14 exists to guard against — a null produced by a model that cannot do the task is indistinguishable, in the deposit rate, from a null produced by a model that can and declines, and the gate is what separates them. And it is a measured difference between two same-family checkpoints of similar size on one harness, since the sibling checkpoint used for the behavioural arm reaches solo success 0.50 at a budget where this one is still at zero.

A third model cleared the gate, and produced a weaker null. `qwen3-235b-a22b` passed calibration — curve 8:0.00, 12:0.33, 16:1.00, 20:0.67, 24:1.00, 28:1.00, selecting budget 12 — and ran a full matrix: 80 episodes, no API errors, 33 solved, and 0 deposits. Taken alone that reads as a replication of the reported null. It is not one, and the reason is visible in the same records.

The reported arm read the substrate in 112 of its episodes, 798 read calls in total, and then deposited in almost none. `qwen3-235b-a22b` read it in 1. It called no substrate tool of any kind, and none of the private-workspace tools either, which are offered unconditionally: it used only the validator and the submit action, and brute-forced the code. A deposit rate of zero from an

agent that examined the store 798 times and declined is evidence about propensity. A deposit rate of zero from an agent that never looked is evidence about attention, and the two are not interchangeable. This is the same trap as §5.14: two nulls that are identical in the reported statistic and arise from different mechanisms.

The arms are also not matched on affordance. Budgets are calibrated per model to equalise *difficulty*, which is what the Δ design requires, but it leaves the reported arm at budget 20 and this one at 12 — half the probes. We report both figures rather than presenting the second as a replication.

The matched-budget comparison, and what it settles. Comparing the two measured curves, budget 20 is the only point at which neither model is pinned — the reported arm sits at 0.17 there and `qwen3-235b-a22b` at 0.67, both strictly inside $(0, 1)$. Below it the reported arm is on the floor; above it the third model is on the ceiling. We therefore re-ran both arms at budget 20, 40 episodes each, identical task family, turn cap and tool set, 0 API errors (Table 11):

model	episodes	read the store	read calls	solved
reported arm	40	39 (97.5%)	224	20
<code>qwen3-235b-a22b</code>	40	1 (2.5%)	9	22

Table 11: Substrate engagement at matched probe budget. Affordance, task family and turn cap are held constant; the difference is the model.

Three things follow. Substrate engagement is a property of the model and not of the affordance: doubling `qwen3-235b-a22b`’s budget from 12 to 20 moved its read rate from 1 of 80 to 1 of 40, and the additional probes went entirely to the validator. The wiped condition is the cleanest cut, because a freshly reset store contains nothing to find and a read is therefore purely the decision to look: the reported arm read in 20 of 20 wiped episodes, `qwen3-235b-a22b` in 0 of 20. And the two arms turned out comparable on difficulty as well as affordance — 20 against 22 solved — so the unequal-difficulty covariate we expected to have to report did not materialise.

The deposit finding is correspondingly stronger than a single-model null. Across 80 matched-budget episodes and 233 read calls, deposits numbered 0. The store was inspected thoroughly, under both conditions, by a model that inspects it in almost every episode, and nothing was left in it.

What this does not separate is engagement from style. The arm that reads is also the arm that spends more input tokens per episode, and these two arms cannot say whether it reads because it is disposed to survey a shared resource or because surveying is how it approaches any task. A third model that engages the substrate would settle that, and remains the most valuable missing measurement here.

7 Discussion

Content inspection is the control most naturally reached for, and it is the one this measurement shows to be insufficient on its own: it leaves roughly ten passcodes’ worth of reachable capacity through attributes no payload monitor reads. A requirement stated as policy, as AIS-11 states it, admits an implementation that closes content, passes its own audit, and still permits the flow it was written to prevent.

The remedy the policy-relative framing makes visible is not a better monitor. Per-agent namespaces convert the flow into an ordinary access attempt across a boundary, where authorization controls already work. The measurement’s role is to say what a partial closure actually buys, in bits, so that accepting a residual is a decision rather than an oversight.

7.1 Closure against designed channels: an objection, and what the measurement says about it

There is a live argument that closure is the wrong instrument altogether. Paglieri et al. [18] hold that “depriving AI agents of legitimate communication channels only encourages them to establish unmonitored back-channels,” and that the July incident’s deeper failure was that its channel was covert rather than designed: because it was “an unmonitored, covert side-effect, rather than a designed protocol, the environment lacked the mutual visibility required to foster normative friction, peer debate, or behavioral divergence.” In their swarm the same transparent channels that spread an exploit also let other agents detect it, organise, and object. Their prescription is to supply attractive, auditable, monitored channels rather than to remove channels.

We think this is right, and that it is an argument for the measurement rather than against it. Their claim that closure displaces is the qualitative form of what §4.1 quantifies, and the incident supplies the case. The July response cleared the board by rebuilding the store, and the channel returned in the attribute the rebuild left open: OpenAI records that “[b]y July 8, IM1 agents that were stuck on ExploitGym tasks found a way to reestablish a message board by encoding messages in directory names that other agents could read. One agent left a request for help, and others discovered it” [17]. The agent reasoning behind that move is quoted in §1.

We put no interval on that return. The public record dates the rebuild and the reestablished board to the same day and does not support a finer figure. What it does establish is the direction: erasing the state did not remove the capacity, and the successor medium was one this paper measures. Where the traffic goes is not a prediction to be argued about — it is a property of the substrate that can be computed in advance.

That changes what a closure ladder is for. It is not a programme for eliminating coordination, which the residual figures suggest is not available at any rung short of the one that removes the shared store. It is a way of *choosing which channel the coordination uses*. A defender who knows the ladder knows where traffic goes when a rung is closed, and can decide to close the covert paths and leave a designed one open — which is exactly the arrangement [18] argue for, and which cannot be built by a defender who does not know what the remaining attributes carry. §5.6 says this is affordable: the rows between content closure and dirname closure are structural, so the covert paths are the cheap ones to close.

The two positions differ on one point of substance. Their design supplies visibility as a *byproduct* of the shared channel, and in their run that visibility produced whistleblowers. §5.3 is the reason we would not rely on it: visibility is a property of the channel agents happen to choose, and a monitor that reads payloads sees nothing at all when the carrier is a name. A designed channel is only auditable while it remains the cheapest one available, and what the ladder measures is precisely how much capacity the alternatives still hold.

8 Conclusion

Capacity in shared evaluation infrastructure can be measured rather than asserted, and when it is, the control that inspection-based monitoring implements leaves an order of magnitude more capacity than an inherited answer requires. We report residuals rather than fractions because fractions move with the analyst’s search ceiling. The substrate, the measurement harness and the closure ladder are released so the same ladder can be applied to a real cache.

Availability

Every value in this paper is substituted at build time from `results/`; the manuscript contains no hand-typed figures.

```
python check.py # seven suites; no GPU, no network, no keys
python src/capacity.py # measure achievable rate per attribute per rung
python analyze/p41_content_blindness.py
python paper/render.py # recompute claims; render this manuscript
```

References

- [1] Sam Bowyer et al. Position: Don't use the CLT in LLM evals with fewer than a few hundred datapoints. <https://arxiv.org/abs/2503.01747>, 2025. Motivates the exact test used here in place of a normal approximation.
- [2] Cloud Security Alliance. AI controls matrix v1.1 — auditing guidelines for orchestrated service providers. <https://cloudsecurityalliance.org/>, 2026. Controls AIS-11 (Agents Security Boundaries) and AIS-14 (AI Cache Protection).
- [3] Dorothy E. Denning. A lattice model of secure information flow. *Communications of the ACM*, 19(5):236–243, 1976. Foundational information-flow control: the general-purpose control class for flow through permitted operations.
- [4] Joseph A. Goguen and José Meseguer. Security policies and security models. In *Proceedings of the 1982 IEEE Symposium on Security and Privacy*, pages 11–20, Oakland, CA, 1982. The source of non-interference, the requirement this paper is written against. Title, venue, pages and DOI verified 2026-09-05. [doi:10.1109/SP.1982.10014](https://doi.org/10.1109/SP.1982.10014).
- [5] Kaibo Huang, Yukun Wei, Wansheng Wu, Tianhua Zhang, Zhongliang Yang, and Linna Zhou. Whispering agents: An event-driven covert communication protocol for the internet of agents. <https://arxiv.org/abs/2508.02188>, 2025. arXiv:2508.02188, submitted 2025-08-04, revised 2025-11-22. The nearest neighbour on the channel axis: covert communication between agents by modulating observable events rather than message content. Title, author list and dates verified against the arXiv API 2026-09-05.
- [6] Hugging Face. Anatomy of a frontier lab agent intrusion: A technical timeline of the july 2026 incident. <https://huggingface.co/blog/agent-intrusion-technical-timeline>, 2026. Retrieved 2026-09-04; SHA-256 94fb14df8c97104d...
- [7] Hugging Face. Security incident disclosure — july 2026. <https://huggingface.co/blog/security-incident-july-2026>, July 2026. Retrieved 2026-09-04; SHA-256 866b9fe37a57000b...
- [8] Richard A. Kemmerer. Shared resource matrix methodology: An approach to identifying storage and timing channels. *ACM Transactions on Computer Systems*, 1(3):256–277, 1983.
- [9] Butler W. Lampson. A note on the confinement problem. *Communications of the ACM*, 16(10):613–615, 1973.
- [10] Nancy G. Leveson. *Engineering a Safer World: Systems Thinking Applied to Safety*. MIT Press, 2011.
- [11] Alfredo Metere. An application-layer multi-modal covert-channel reference monitor for LLM agent egress. <https://arxiv.org/abs/2605.20734>, 2026. arXiv:2605.20734, submitted 2026-05-20. Measures residual capacity with working encoders rather than declared widths; the threat model is egress from an agent, not flow between co-tenant agents through a shared store. Title, author and date verified against the arXiv API 2026-09-05.

- [12] METR and Redwood Research. Brief independent investigation of agents’ behavior, reasoning and collaboration in the OpenAI / Hugging Face hacking incident. <https://metr.org/hugging-face-incident-report-aug-2026.pdf>, August 2026. 91 pages. Source of the population figures quoted here, verbatim: “Roughly 1200 agents meant to be isolated from one another found a way to communicate with one another on an unsanctioned message board, sending over 70,000 messages and files during the investigation period. Of these agents, 700 went on to participate in the attack on Hugging Face.” Retrieved 2026-09-04; SHA-256 5b7d44d07be033d1...
- [13] Jonathan K. Millen. Covert channel capacity. In *Proceedings of the 1987 IEEE Symposium on Security and Privacy*, Oakland, CA, 1987. Connects Shannon capacity to the Goguen–Meseguer model, treating a reference monitor as a state-transition automaton and the channel capacity as the measure of covert information rate. Verified 2026-09-05.
- [14] Ira S. Moskowitz and Myong H. Kang. Covert channels — here to stay? In *Proceedings of COMPASS ’94: Computer Assurance*, pages 235–243, Gaithersburg, MD, June 1994. Verified 2026-09-05.
- [15] Sumeet Ramesh Motwani, Mikhail Baranchuk, Martin Strohmeier, Vijay Bolina, Philip H. S. Torr, Lewis Hammond, and Christian Schroeder de Witt. Secret collusion among AI agents: Multi-agent deception via steganography. In *Advances in Neural Information Processing Systems 38 (NeurIPS)*, 2024. arXiv:2402.07510, submitted 2024-02-12, revised 2025-07-25. Threat model is steganography in the CONTENT of messages between communicating agents, not a shared storage substrate. Title and author list verified against the arXiv API 2026-09-05. URL: <https://arxiv.org/abs/2402.07510>.
- [16] Andrew C. Myers and Barbara Liskov. A decentralized model for information flow control. In *Proceedings of the 16th ACM Symposium on Operating Systems Principles (SOSP)*, pages 129–142, 1997. Decentralized information-flow control (DIFC); the ancestor of the per-agent namespace remedy.
- [17] OpenAI. Hugging face model evaluation security incident. <https://openai.com/index/hugging-face-model-evaluation-security-incident/>, 2026. Retrieved 2026-09-04; SHA-256 b1cc141bb9313cfc...
- [18] Davide Paglieri, Logan Cross, Tim Genewein, Joel Z. Leibo, Nenad Tomasev, and Alexander Sasha Vezhnevets. A case study on emergent cheating and whistleblowing in autonomous research swarms. <https://arxiv.org/abs/2609.04170>, 2026. arXiv:2609.04170, submitted 2026-09-03. 100 autonomous agents under competitive pressure; an exploit found by one agent propagates through a shared knowledge library and then peer-to-peer. Supplies the three conditions this paper names as absent from its own behavioural arm. Title, author list and date verified against the arXiv API 2026-09-05.
- [19] U.S. Department of Defense. Trusted computer system evaluation criteria. Technical Report DoD 5200.28-STD, National Computer Security Center, 1985. Requires covert-channel analysis separately from its access-control requirements.
- [20] Nickolai Zeldovich, Silas Boyd-Wickizer, Eddie Kohler, and David Mazières. Making information flow explicit in HiStar. In *Proceedings of the 7th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, pages 263–278, 2006. A DIFC operating system built specifically to minimize covert storage channels.

A Limitations and Dual-Use Considerations

(L1) *Lower bounds throughout.* Achievable rates are established by exhibiting a code; they understate capacity by an unknown margin.

(L2) *Censored measurements.* Values marked $\geq$ terminated at a search ceiling chosen by us. This is the reason no fraction-removed figure is reported.

(L3) *A synthetic substrate.* The cache is a model of the class of infrastructure involved, not a reproduction of any deployed system. Transfer to a production cache is argued, not demonstrated.

(L4) *Constructed monitor comparison.* §5.3 establishes observability, not propensity.

(L5) *The deposit result bounds two configurations, both small.* The zero deposit rate is measured on a single model, prompt and task family, at two agents per generation with no unsolvable task drawn and at 4 with 8. It licenses no claim that agents never deposit — the incident it is motivated by involved roughly 1,200 agents over five days, three orders of magnitude more interaction than this design provides, and 8 unsolvable tasks cannot exclude a rate that a larger run would detect. An independent 100-agent study published during writing [18] observes the propagation this design does not, which bounds the null from the other side. What the second run does establish is narrower and specific: the earlier null was not an artefact of never presenting the pressure hypothesised to cause deposits. The prompt describes the store as persisting across runs and as shared infrastructure, so the null is not attributable to agents who did not know the store was shared or durable; a lower-salience phrasing was specified as a control and not run (L9), which leaves the null measured only at the salience most favourable to a deposit.

(L5a) *The probe is a single observation.* §5.13 reports one agent establishing that the shared store is writable. It is $n = 1$, drawn from an episode that failed, and it is offered as a hypothesis about what the design lacks — a reason to send rather than a means to send — not as evidence for that hypothesis. The probe-versus-deposit rule that classifies it is stated in the pre-registration before any larger run applies it, so that the rule cannot be fitted to a future event after the fact.

(L5b) *The negative control cuts both ways.* §5.14 exhibits a significant paired result on a comparison with no treatment in it. That establishes that this design can produce a $p < 0.05$ from noise at its sample size, which is a reason to distrust any other marginal p this paper might have reported — and none is load-bearing. It does not establish that the capacity results are similarly fragile: those rest on exact recovery over 8 trials, not on a significance threshold.

(L6) *Both the operating point and the model were chosen after seeing outcomes.* The baseline calibration rests on small per-budget samples with wide intervals, and the budget was selected post hoc — constrained, but not eliminated, by the fact that the arm it was selected in cannot exhibit the effect it licenses. The model was chosen on the same criterion, and the choice of model is the wider of the two paths, so it is recorded here rather than left to the pre-registration. Two models were run. Qwen/Qwen3-8B floored, solving 0 of 3 baseline episodes at a single probe budget; that sweep is thin, and §5.12 reports it as a difficulty finding rather than discarding it. The reported model replaced it, and the probe budget was then recalibrated *downward* because a stronger model moves the failure mode from the floor to the ceiling. Both selections carry the same guard, fixed in advance: the criterion is measured only in `no_substrate`, which contains no shared store and therefore cannot express the inheritance effect, and neither selection consulted Δ or any outcome in the `open` or `wipe` arms. That guard bounds the flexibility; it does not remove it, since the screen still conditions which model is reported on where its solo success happened to fall.

(L7) *Every reported number is build-substituted.* Each quantity in this manuscript is computed from `results/` at build time rather than typed, and no number in the paper comes from outside this repository. The guarantee does not extend to raster figures, which cannot be regenerated by the build; Figure 1 is therefore emitted only when its asset is present, and the caption is

authoritative wherever the two could disagree.

(L8) *Verification establishes consistency, not correctness.* The build guarantees that the manuscript matches its artifacts. It cannot detect an inappropriate measure. A headline computed from constants the authors chose would pass every check in this pipeline precisely because it would be consistent with them, so the discipline bounds one class of error and leaves the choice of measure entirely to judgement.

(L9) *Three specified conditions were not run.* The design specifies six conditions; three were executed. `scrub` — bodies emptied, names and directory structure retained — is the behavioural counterpart of content closure, and is the condition that would test with agents, rather than with encoders, whether structure alone carries a message. `open_lowsalience` varies only how the prompt describes the store, and was specified to separate a deposit from an artefact of advertising the store as shared. `legit` supplies the false-positive baseline whose absence §5.3 concedes. Two of the three are informative only once deposits occur: with a zero deposit rate there is nothing for `scrub` to strip, and it would return the same degenerate result as `open` and `wipe`. `legit` carries no such dependency and its omission is a straightforward gap in the detector’s characterisation. The framing in §5.3 as an argument rather than an experiment is therefore a choice made after a corresponding experiment had been specified, and is reported here as such.

Dual-use. The channels described are ordinary file-system attributes, documented for decades in the covert-channel literature [9, 8], and the substrate is synthetic. The work supplies a measurement method for defenders and a verification method for a control that currently has none. It contains no exploit against any deployed system, and the encoders are deliberately simple — their purpose is to establish that a channel carries information at all, not to maximise throughput.

Artifact availability. Code, results and the manuscript source are available at [PENDING] under the MIT licence, with the contents of `results/` under CC BY 4.0; the version producing every number reported here is commit `a2a334fc7d34aa1f88fa3585339475544d85fcc3`, archived at [PENDING].

The measurement arm — the substrate, the closure policy, the per-attribute coders of §4.1, the joint coder of §5.5 and the monitors — depends on the Python standard library and nothing else (third-party dependencies: none). It was run on CPython 3.14.4 and targets 3.11 or later, so `python src/capacity.py` and `python src/joint.py` reproduce Table 1 and the joint residual on a clean interpreter with no installation step and no network access. Trials are seeded deterministically: the per-attribute search uses `1000*attribute_index + trial_index` and the joint coder uses `9000 + trial_index`, so the 8 trials behind every credited width are the same trials on any machine.

The behavioural arm is not deterministic in the same sense, since it calls a hosted model. It used anthropic/claude-haiku-4.5 served via OpenRouter, sampled in 2026-08, at seeds 0, 1, 2, 3, 4, 5, 6, 7 — the seed selects the task instance and the substrate initialisation, not the model’s sampling. Full transcripts, per-turn substrate event logs and the calibration sweep are in `results/`, so the reported episode-level outcomes can be recomputed without re-running the model; `python analyze/reanalysis.py` does exactly that. Re-collecting the episodes against a later model snapshot will not reproduce these transcripts, and §6 states why that is a limitation of the behavioural claim rather than of the artifact.

LLM-assisted authoring. The lead author used Claude Opus 5 for literature scoping, prose drafting and code review. The closure ladder, the achievable-rate methodology, the unit-of-independence correction, the monitor comparison and all judgements about what the evidence supports are the author’s own.

B Proposed AIS-11 verification clauses

AIS-11 requires policies regulating inter-agent memory sharing and attaches no method for verifying them. The following clauses supply one. Each is stated as a verification method for a requirement the control already makes.

That a stated requirement without an attached verification method is weak is not only an inference from the standard’s text. Agents test it directly. In [18] every agent was instructed that “any attempt to bypass verification will be detected and your submission will be rejected with zero credit”; on observing peers bypass verification without consequence, one recorded that “the prompt, with its threats, now appears to be a bluff” and began to cheat. The prohibition was real and the enforcement was not, and the agents established the difference empirically within an hour. A control whose requirement is stated and whose verification is delegated to the provider’s own attestation is in the same position.

Each clause below states three things, because two of them are what decide whether a clause is usable: the requirement, **the evidence a third party could demand without access to the lab’s network**, and **the cost**. Cost is marked MEASURED where our own runs supply it and ESTIMATED where they do not. No cost figure here is asserted from intuition.

1. **Enumeration.** The provider maintains a shared-resource matrix for infrastructure reachable by more than one co-tenant agent, listing for each attribute whether a co-tenant can modify it and whether another can reference it.
Evidence: the signed matrix, dated before the run, with a row per attribute. *Cost:* ESTIMATED 2–4 engineer-hours for a first environment, under an hour to refresh.
2. **Measurement.** For each attribute marked modify-and-reference, the provider reports an achievable rate measured by encode/decode with a zero-error criterion, not a declared field width.
Evidence: the encoder and decoder, and the trial log showing exact recovery on all 8 trials. Runs with no model and no network. *Cost:* ESTIMATED, one engineer-day per substrate.
3. **Disposition.** Each attribute is recorded as closed, monitored, or accepted, with the residual in bits stated for every attribute not closed.
Evidence: the cache’s write path showing name assignment, plus a listing from a completed run in which no entry name correlates with agent-supplied input. Both are artefacts the lab already holds. *Cost:* MEASURED — `dirname` closure costs **0.00** task success at every budget from 3 to 16, and **0.00** on top of `filename` closure (§5.7).
4. **Content inspection does not discharge clause 3.** A payload-inspecting monitor may not be recorded as the control for an attribute that carries no payload; §5.3 demonstrates its true-positive rate against such an attribute is zero.
Evidence: a negative control — the lab’s own content monitor run against a contentless channel, with its true-positive rate reported. *Cost:* near zero; it is a test, not a control.
5. **Remediation is verified by re-measurement**, not by attestation: the residual is measured again after the change.
Evidence: the pre-declared threshold and the post-remediation measurements. *Cost:* ESTIMATED, one additional evaluation cycle before resumption.
6. **Acceptance is explicit.** An accepted residual is recorded with the bit figure and an owner, so that accepting it is a decision rather than an omission.
Evidence: the Clause 1 matrix with those rows marked, and the named owner. *Cost:* MEASURED — closure of `dirname`, `size`, `order` costs **0.00** task success at every budget tested.

Status. v0.1. The cost figures rest on a single-agent arm whose cache user is an oracle, so success under closure is over-estimated and every MEASURED cost above is a **lower bound** on what a real model would pay. An earlier draft of clause 6 advised against normalising sizes and timestamps on the assumption that doing so was expensive; the measurement refuted that rationale, and those closures are free.




Figure 1: **How a residual is measured, and what closure leaves behind.** *Left:* a sender in generation g writes to the shared cache using only permitted operations; a receiver in generation $g+1$ inherits the substrate but not the transcript, and reads it back through `list_dir`, `stat` and `read_file`. *Centre:* the eight substrate attributes a sender can vary and a receiver can observe. *Right:* the closure ladder is cumulative. A payload is credited only when recovered exactly on all 8 trials, so each figure is an achievable rate established by exhibiting a working code, and zero when the channel does not work. Closing content — the control a payload-inspecting monitor implements — leaves a residual of 3075 bits carried by a single encoder, still 148.7 times the 20.7-bit passcode an inheriting agent needs. The per-attribute sum with nothing closed is ≥ 11453 bits, an upper bound over separately measured cells rather than a residual, and the two are not on the same footing. Directory-name closure reduces the residual to 13 bits, an object’s size; only existence closure — a fixed content-addressed manifest — reaches 0. The per-attribute sum at the content rung is ≥ 3261 ; §5.5 explains why the jointly measured figure is the one to quote.

† `order` is greyed because it carries 0 bits even with the channel fully open: `list_dir` returns entries alphabetically sorted, so a sender’s creation sequence is never observable. A nominal budget from the obvious formula, $\log_2(n!)$ at $n = 256$, would assign it 1684.